



CORNER DETECTION

Intuition of a Corner




- Intuitively they are “junctions of contours”.

Why do we care about corners?




- They can help us to match two images.
- Useful to recover 3D structure, create panoramic images, tracking, and more!
- Good features to match!


Harris Corner Detection



“flat” region:
no change in
all directions



“edge”:
no change along
the edge direction



“corner”:
significant change
in all directions

Harris corner detector gives a mathematical approach for determining which case holds.

Harris Corner Detection (cont'd)

$$E(u, v) = \sum_{x, y} w(x, y) (I(x + u, y + v) - I(x, y))^2$$

The diagram shows the equation $E(u, v) = \sum_{x, y} w(x, y) (I(x + u, y + v) - I(x, y))^2$. A red oval highlights the term $(I(x + u, y + v) - I(x, y))^2$. Three blue arrows point from the terms in the equation to labels below: from $w(x, y)$ to "Window (e.g., Gaussian)", from $I(x + u, y + v)$ to "Shifted Intensity", and from $I(x, y)$ to "Intensity". A red arrow points from the red oval to the text below.

Window
(e.g., Gaussian)

Shifted Intensity

Intensity

- For nearly constant patches, the difference is small.
- For distinctive patches, the difference is large.

Taylor series for 2D functions

- First order approximation:

$$f(x + u, y + v) \approx f(x, y) + u f_x(x, y) + v f_y(x, y)$$

Harris Corner Detection (cont'd)

Taylor approximation



$$\begin{aligned}\sum_{x,y} (I(x+u, y+v) - I(x, y))^2 &\approx \sum_{x,y} (I(x, y) + uI_u(x, y) + vI_v(x, y) - I(x, y))^2 \\&= \sum_{x,y} u^2 I_u^2 + 2uv + v^2 I_v^2 \\&= \sum_{x,y} \begin{bmatrix} u & v \end{bmatrix} \begin{bmatrix} I_u^2 & I_u I_v \\ I_v I_u & I_v^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} \\&= \begin{bmatrix} u & v \end{bmatrix} \left(\sum_{x,y} \begin{bmatrix} I_u^2 & I_u I_v \\ I_v I_u & I_v^2 \end{bmatrix} \right) \begin{bmatrix} u \\ v \end{bmatrix}\end{aligned}$$

Harris Corner Detection (cont'd)

- Rewriting

$$E(u, v) \approx \begin{bmatrix} u & v \end{bmatrix} M \begin{bmatrix} u \\ v \end{bmatrix}$$

- The matrix M (2x2) is




$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_u^2 & I_u I_v \\ I_v I_u & I_v^2 \end{bmatrix}$$

Intuition about Harris

- Assume gradients are 2D points (I_x, I_y)
- Fit an ellipse to the set of points via scatter matrix
- Analyze ellipse parameters to determine if we detected a corner

Intuition about Harris (visualization)

The distribution of x and y derivatives can be characterized by the shape and size of the principal component ellipse



Classification via Eigenvalues

- If both eigenvalues are large \Rightarrow corner!
- If single eigenvalue is large \Rightarrow edge.
- If both eigenvalues are small \Rightarrow flat region
- Computing eigenvalue is expensive!

Measure of corner response


- Corner response:

$$R = \det M - k \operatorname{trace}(M)^2$$

- Constant k in the range (0.02 – 0.06)
- Determinant is the product of the eigenvalues
- Trace is the sum of the eigenvalues

Corner response map

- R depends only on eigenvalues of M
- R is large for a corner
- R is negative with large magnitude for an edge
- $|R|$ is small for a flat region



Usually $R > 10000$

Harris corner detector algorithm

1. Compute x and y derivatives of image

$$I_x = G_\sigma^x * I \quad I_y = G_\sigma^y * I$$

2. Compute products of derivatives at every pixel

$$I_{x2} = I_x \cdot I_x \quad I_{y2} = I_y \cdot I_y \quad I_{xy} = I_x \cdot I_y$$

3. Compute the sums of the products of derivatives at each pixel

$$S_{x2} = G_{\sigma'} * I_{x2} \quad S_{y2} = G_{\sigma'} * I_{y2} \quad S_{xy} = G_{\sigma'} * I_{xy}$$

4. Define at each pixel (x, y) the matrix

$$H(x, y) = \begin{bmatrix} S_{x2}(x, y) & S_{xy}(x, y) \\ S_{xy}(x, y) & S_{y2}(x, y) \end{bmatrix}$$

5. Compute the response of the detector at each pixel

$$R = \text{Det}(H) - k(\text{Trace}(H))^2$$

6. Threshold on value of R . Compute nonmax suppression.

Shi-Tomasi corner test


$$\min(\lambda_1, \lambda_2) > \lambda$$



Threshold


- Accept as corner if above threshold.

FAST corners: Features from Accelerated Segment Test



- Uses a Bresenham circle of radius 3 (16 pixels).
- Defines a fast test to classify pixels as corners or non corners.


FAST corners: Features from Accelerated Segment Test



Condition:

1. If N (typically 12) contiguous pixels are brighter/darker than the current intensity \pm threshold, then classify as corner.

High-speed test of non-corners



- Examine pixels 1, 5, 9, 13.
- If current pixel is a corner, then at least 3 pixels examinations must be brighter/darker than $I_p + t$ / $I_p - t$, respectively.
- If this is not satisfied, reject.

Weaknesses of the high-speed test.

1. The high-speed test does not generalize well for $n < 12$.
2. The choice and ordering of the fast test pixels contains implicit assumptions about the distribution of feature appearance.
3. Knowledge from the first 4 tests is discarded.
4. Multiple features are detected adjacent to one another.

Using machine learning to address the weaknesses

- Detect corners to generate a training set.
- For every pixel I_x in the circle (1-16) of the training set, label them as:
 - Darker if $I_x \leq I_p - t$
 - Similar if $I_p - t \leq I_x \leq I_p + t$
 - Brighter if $I_x \geq I_p + t$
- Train a decision tree for every pixel in the circle.
- The decision tree can be implemented efficiently in C/C++.

BLOB DETECTION

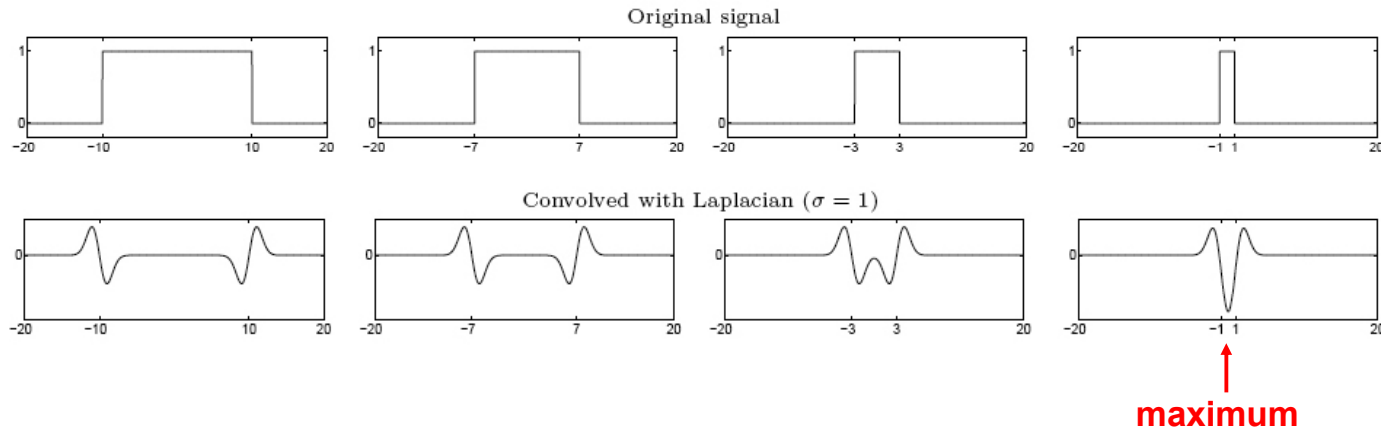
Intuition of a blob



A “radially” contrasting region


Blobs in 1D

- Edge = ripple
- Blob = superposition of two ripples



- The magnitude of the Laplacian response will achieve a maximum at the center of the blob.
- This happens when the scale of the Laplacian filter (i.e., its sigma parameter) matches the scale of the blob.
- Another way of thinking about this is when the similarity, measured by convolution, between the signal and the Laplacian is maximum.

Laplacian response decays as scale increases




Scale normalization






- To keep the responses the same (scale-invariant), we must multiply Gaussian derivative by σ .
- Laplacian is the second derivative of a Gaussian, so we must multiply by σ^2 .

Effect on scale normalization






Original signal



Unnormalized Laplacian response





Scale-normalized Laplacian response



↑
maximum

Blob detection in 2D



Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D



$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

Blob detection in 2D normalized

Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D




Scale-normalized:
$$\nabla_{\text{norm}}^2 g = \sigma^2 \left(\frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$

Scale selection


- What is the σ parameter of the LoG so that we can detect a circle of radius r ?
- To get maximum response, the zeros of the LoG have to be aligned with the circle. This happens at

$$\sigma = \frac{r}{\sqrt{2}}$$

Scale selection




image



Characteristic scale of a blob


- The characteristic scale of a blob is the one that produces the maximum LoG response.



T. Lindeberg (1998). ["Feature detection with automatic scale selection."](#)
International Journal of Computer Vision **30** (2): pp 77--116.

Overview of Scale-space blob detector

1. Convolve the image at several scales with scale-normalized LoG.
2. Find the the maximal response in scale-space.



Why do you think corners or blobs are used more frequently than edges?

- Corners and blobs are repeatable
- Corners and blobs can be localized more easily